

Neural Processes on Graphs

Huidong (Leon) Liang

Discipline of Business Analytics
The University of Sydney Business School

June 2021

Overview

1 Neural Processes

- Variational Autoencoders
- Meta-learning
- Gaussian Processes
- Neural Processes

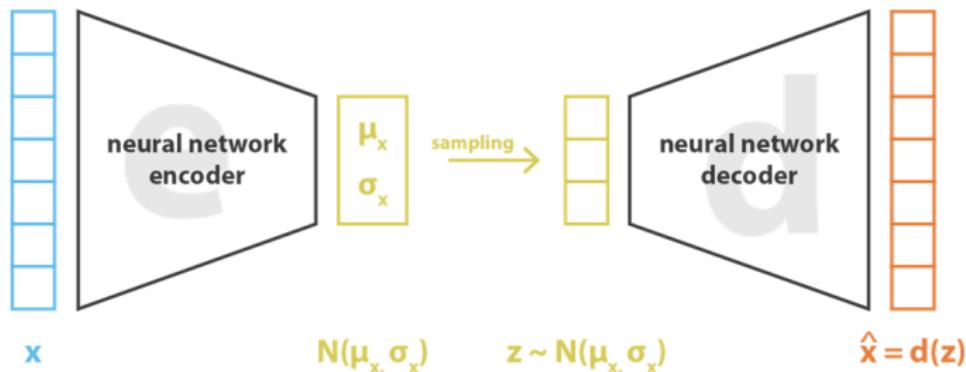
2 Neural Processes on Graphs

- Definition of Network/Graph
- Spectral Graph Convolution
- Variational Graph Autoencoders
- Graph Neural Processes

Neural Processes

Related Work: Variational Autoencoders

Model Structure



$$\mu = \text{ReLU}(W_\mu \text{ReLU}(W_1 x + b_1) + b_\mu) \quad (1)$$

$$\log \sigma = \text{ReLU}(W_\sigma \text{ReLU}(W_1 x + b_1) + b_\sigma) \quad (2)$$

$$\log p(x|z) = \sum_{i=1}^d x_i \log \hat{x}_i + (1 - x_i) \log(1 - \hat{x}_i) \quad (3)$$

$$\text{where } \hat{x} = \sigma(W_3 \sigma(W_2 z + b_2) + b_3) \quad (4)$$

Related Work: Variational Autoencoders

Inference

- Bayesian Inference

$$p(z|x) = \frac{p(x|z)p(z)}{p(x)} \propto p(x|z)p(z) \quad (5)$$

- Variational Inference

We want to use a variational distribution $q(z|x)$ to approximate $p(z|x)$ by minimizing the Kullback-Leibler (KL) Divergence:

$$KL(q(z|x)||p(z|x)) = \mathbb{E}_{q(z|x)} \left[\log \frac{q(z|x)}{p(z|x)} \right]$$

Related Work: Variational Autoencoders

Inference

- We can also express KL Divergence as:

$$\log p(x) = \mathbb{E}_{q(z|x)}[\log p(x)] \quad (6)$$

$$= \mathbb{E}_{q(z|x)} \left[\log \frac{p(x, z)}{p(z|x)} \right] \quad (7)$$

$$= \mathbb{E}_{q(z|x)} \left[\log \left[\frac{p(x, z)}{q(z|x)} \frac{q(z|x)}{p(z|x)} \right] \right] \quad (8)$$

$$= \mathbb{E}_{q(z|x)} \left[\log \frac{p(x, z)}{q(z|x)} \right] + \mathbb{E}_{q(z|x)} \left[\log \frac{q(z|x)}{p(z|x)} \right] \quad (9)$$

- We define $\mathcal{L} = \mathbb{E}_{q(z|x)} \left[\log \frac{p(x, z)}{q(z|x)} \right]$ as the Variational Lower Bound.

Related Work: Variational Autoencoders

Inference

- Now we rearrange the above equations to have a new form for \mathcal{L} :

$$\mathcal{L} = \mathbb{E}_{q(z|x)} \left[\log p(x) \right] - \mathbb{E}_{q(z|x)} \left[\log \frac{q(z|x)}{p(z|x)} \right] \quad (10)$$

$$= \mathbb{E}_{q(z|x)} \left[\log \frac{p(x|z)p(z)}{p(z|x)} \right] - \mathbb{E}_{q(z|x)} \left[\log \frac{q(z|x)}{p(z|x)} \right] \quad (11)$$

$$= \mathbb{E}_{q(z|x)} \left[\log p(x|z) \right] - \mathbb{E}_{q(z|x)} \left[\log \frac{q(z|x)}{p(z)} \right] \quad (12)$$

$$= \mathbb{E}_{q(z|x)} \left[\log p(x|z) \right] - \mathbf{KL}\left(q(z|x)||p(z)\right) \quad (13)$$

Related Work: Variational Autoencoders

Inference

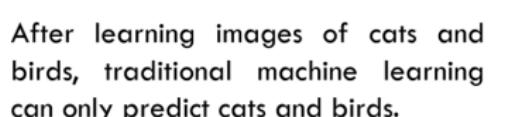
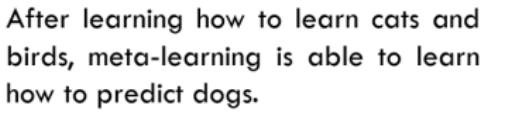
- To optimize \mathcal{L} , we need to calculate the closed-form gradient $\nabla \mathcal{L}$, which requires the *reparameterization trick* when sampling z :

$$\mathbb{E}_{q(z|x)}[\log p(x|z)] \simeq \frac{1}{L} \sum_{l=1}^L \log p(x|z^{(l)}) \quad (14)$$

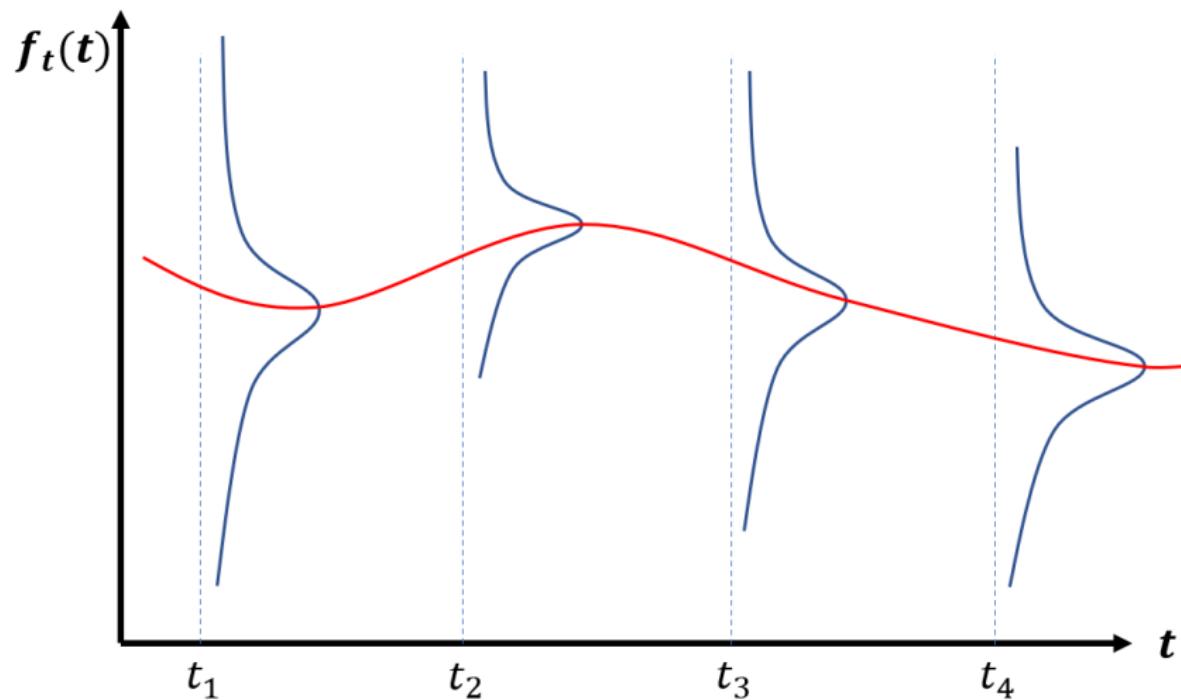
$$z^{(\ell)} = \mu + \sigma \epsilon^{(\ell)}, \text{ with } \epsilon^{(\ell)} \sim \mathcal{N}(0, 1) \quad (15)$$

- Here we use an auxiliary noise variable $\epsilon \sim \mathcal{N}(0, 1)$ and a function $z^\ell = \mu + \sigma \epsilon^{(\ell)}$ to replace $z^\ell \sim \mathcal{N}(\mu, \sigma)$.

Related Work: Meta-learning

	Learning	Predicting
Traditional Machine Learning	 	 
Meta-Learning	 	 

Related Work: Gaussian Processes



$$\{f_t(t)\}_{t \in T} \sim \text{GP}(\mu(t), k(t, t'))$$

Neural Processes

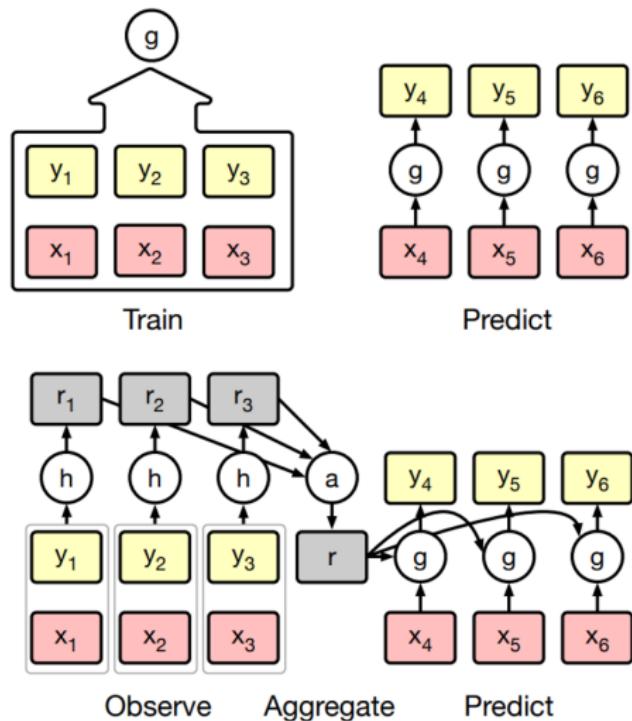
Overall Idea

Conditional Neural Processes

- We have a group of p datasets $\{D_1, \dots, D_p\}$, and we assume each dataset D_i is generated from a function $y = f_i(x)$.
- $\{f(x)\}_p \sim GP(\mu(x), k(x, x'))$
For each $y = f_i(x)$, we sample some x and y pairs.

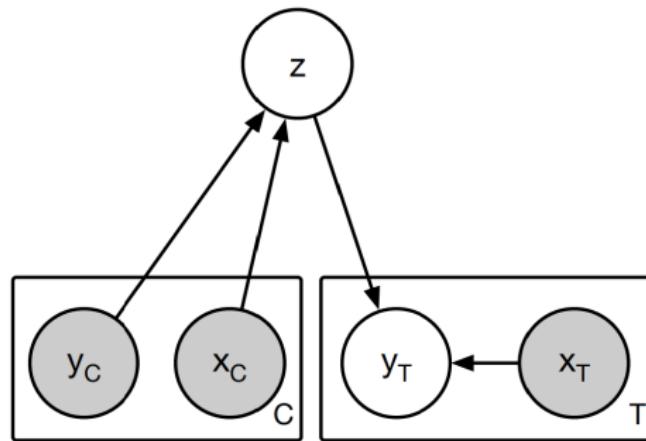
$$\begin{array}{ccccccccc} x_1 & y_1 & x_4 & y_4 & & & & & \\ x_2 & y_2 & x_5 & y_5 & \dots & & x_{49} & y_{49} & \\ x_3 & y_3 & x_6 & y_6 & & & x_{50} & y_{50} & \\ f_1(x) & f_2(x) & \dots & f_p(x) & & & & & \end{array}$$

- Since we cannot distinguish different datasets all the time, we will then treat them as one large dataset, and we split this complete dataset into two sets:
 - Context Set $C = \{x_C, y_C\}_{i=1}^n$
 - Target Set $D = \{x_D, y_D\}_{i=n+1}^{n+m}$



Neural Processes

Model Structure: Encoder



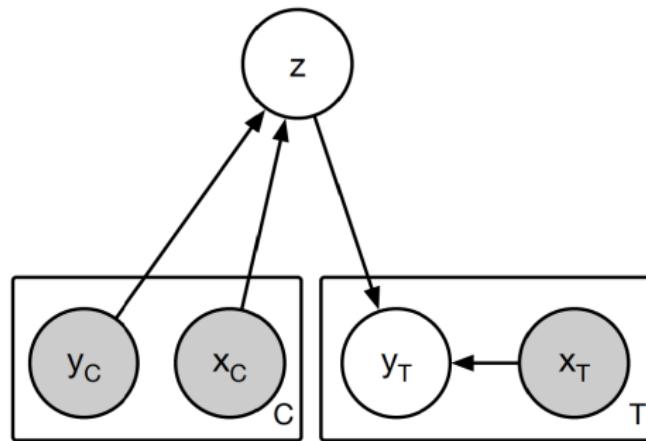
$$\mu_c = \text{ReLU}(W_{\mu_c} \text{ReLU}(W_1[x_C \ y_C] + b_1) + b_\mu) \quad (16)$$

$$\log \sigma_c = \text{ReLU}(W_{\sigma_c} \text{ReLU}(W_1[x_C \ y_C] + b_1) + b_\sigma) \quad (17)$$

- Here μ_c and σ_c^2 contains means and variances of r_c for each data pair $\{x_C, y_C\}_i$ in the Context Set C .

Neural Processes

Model Structure: Aggregation



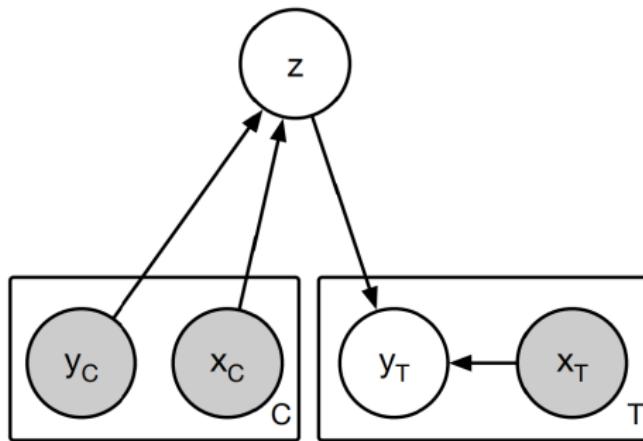
$$\mu_z(r) = \frac{1}{m} \sum_{i=1}^m \mu_{c_i} \quad \log \sigma_z(r) = \frac{1}{m} \sum_{i=1}^m \log \sigma_{c_i} \quad (18)$$

- Use $\mu_z(r)$ and $\sigma_z^2(r)$ to parameterise the distribution for the global latent representation z .

$$q(z|C) = \mathcal{N}(\mu_z(r), \sigma_z^2(r))$$

Neural Processes

Model Structure: Decoder



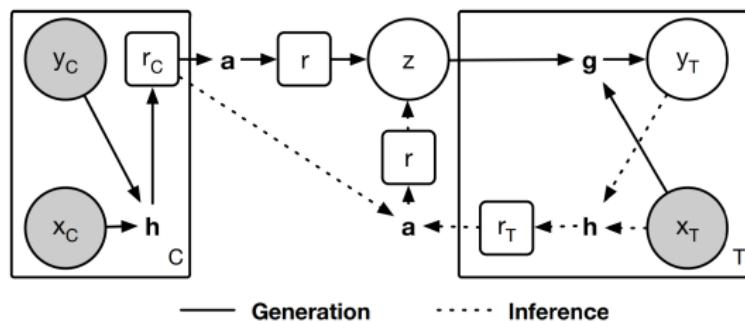
- Combine a sampled z with every feature $x_{T_i} \in X_T$ as $\tilde{x}_{T_i} \in \tilde{X}_T$, where $\tilde{x}_{T_i} = [x_{T_i}^\top z^\top]^\top$.

$$\log p(x|x_T, z) = \sum_{i=1}^d y_{T_i} \log \hat{y}_{T_i} + (1 - y_{T_i}) \log(1 - \hat{y}_{T_i}) \quad (19)$$

$$\text{where } \hat{y}_T = \sigma(W_3 \sigma(W_2 \tilde{x}_{T_i} + b_2) + b_3) \quad (20)$$

Neural Processes

Inference and Learning



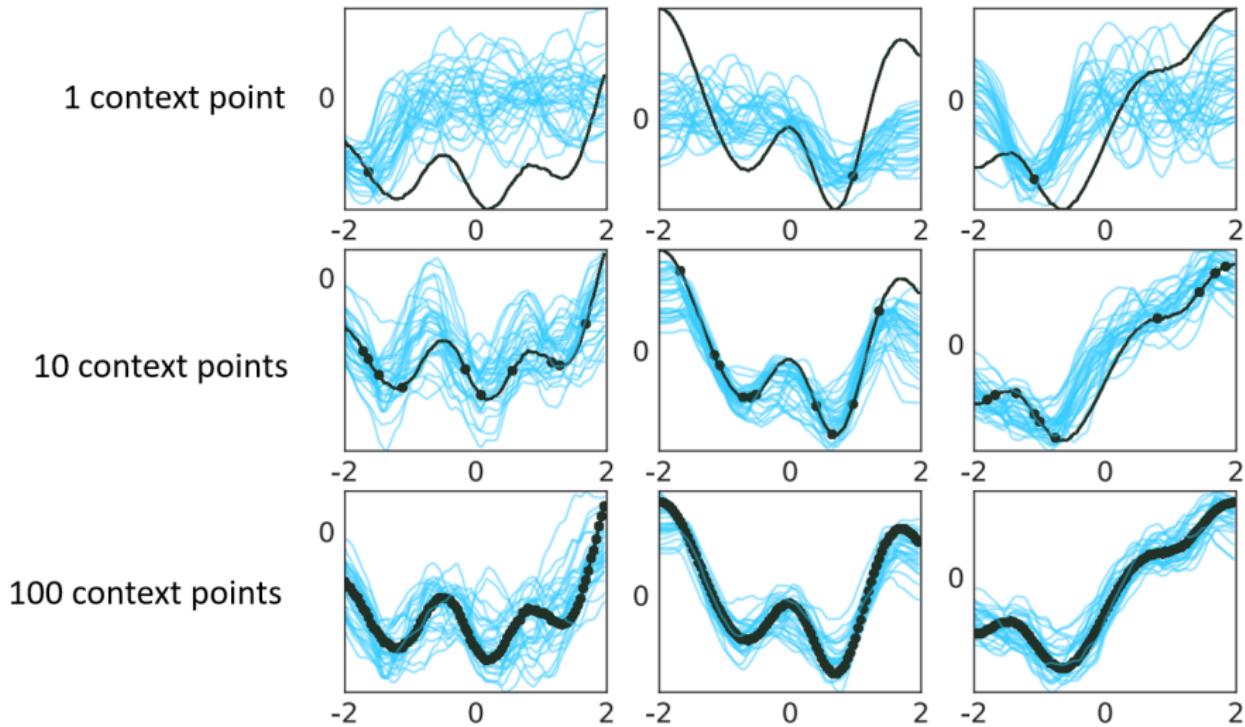
- We want to optimize the variational lower bound \mathcal{L} :

$$\mathcal{L} = \mathbb{E}_{q(z|D)} \left[\log p(y_T|x_T, z) \right] - \mathbf{KL}\left(q(z|D) || q(z|C) \right) \quad (22)$$

Here $D = C \cup T$, where C is our Context set and T is our Target set.

Neural Processes

An Example on Neural Processes



Neural Processes on Graphs

Neural Processes on Graphs

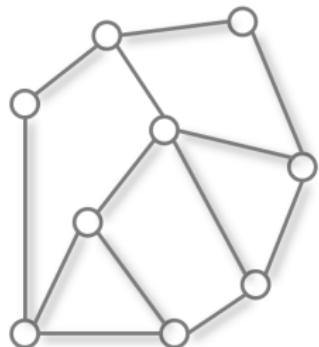
Definition of a Network/Graph



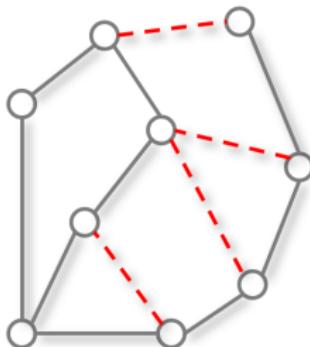
- We define an undirected graph $G = (V, A)$, where V is the set of nodes with features $x_i \in X$ corresponding to each node $v_i \in V$, and the adjacency matrix A with $A_{ij} = 1$ if there is a link between v_i and v_j , and $A_{ij} = 0$ otherwise.

Neural Processes on Graphs

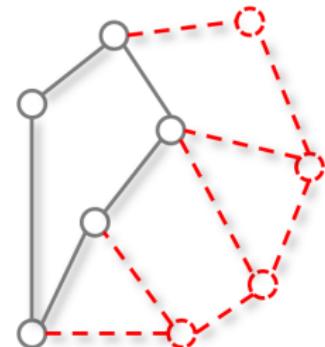
Link Prediction on Network/Graph



Original Graph



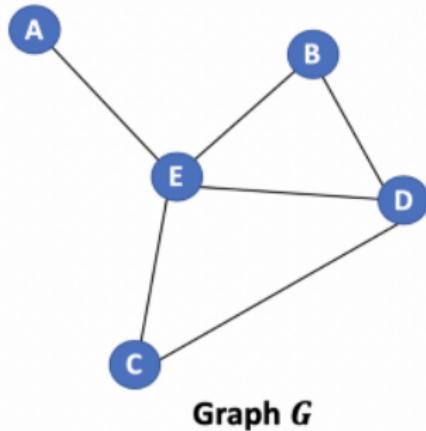
Transductive
Link Prediction



Inductive
Link Prediction

Neural Processes on Graphs

Related Work: Spectral Graph Convolutional Networks (GCN)



	A	B	C	D	E
A	0	0	0	0	1
B	0	0	0	1	1
C	0	0	0	1	1
D	0	1	1	0	1
E	1	1	1	1	0

Adjacency matrix A

	A	B	C	D	E
A	1	0	0	0	0
B	0	2	0	0	0
C	0	0	2	0	0
D	0	0	0	3	0
E	0	0	0	0	4

Degree matrix D

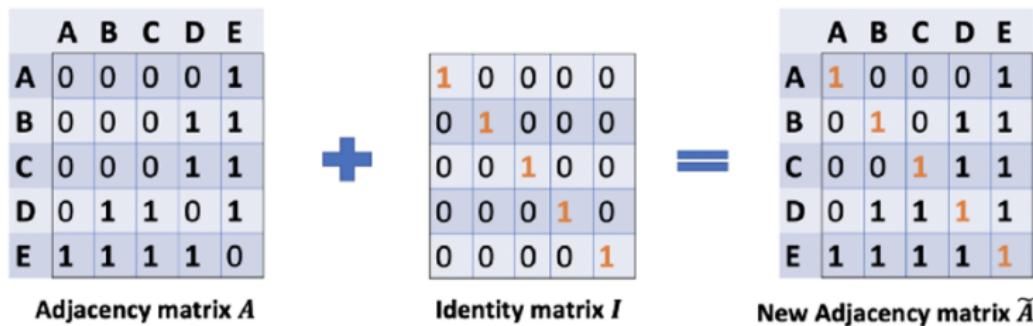
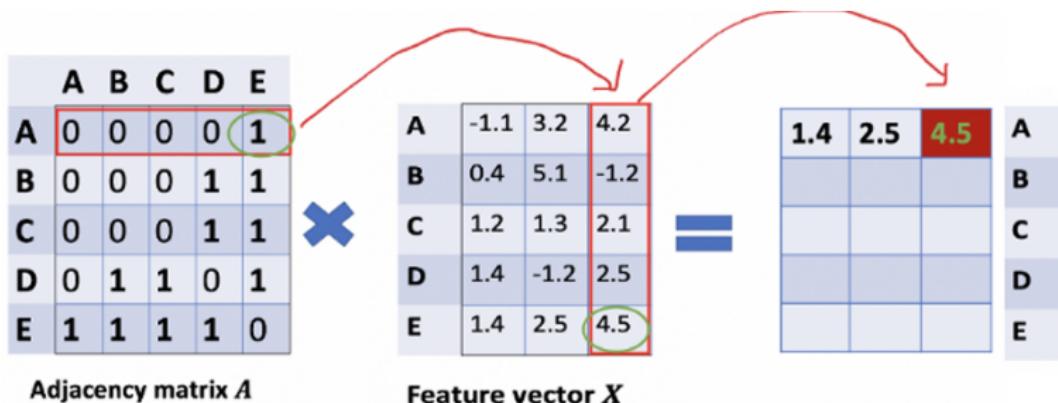
	A	B	C
A	-1.1	3.2	4.2
B	0.4	5.1	-1.2
C	1.2	1.3	2.1
D	1.4	-1.2	2.5
E	1.4	2.5	4.5

Feature vector X

<https://www.topbots.com/graph-convolutional-networks/>

Neural Processes on Graphs

Related Work: Spectral Graph Convolutional Networks (GCN)



Neural Processes on Graphs

Related Work: Spectral Graph Convolutional Networks (GCN)

	A	B	C	D	E
A	1	0	0	0	1
B	0	1	0	1	1
C	0	0	1	1	1
D	0	1	1	1	1
E	1	1	1	1	1

New adjacency matrix \tilde{A}

2	0	0	0	0	0
0	3	0	0	0	0
0	0	3	0	0	0
0	0	0	4	0	0
0	0	0	0	5	

New degree matrix \tilde{D}

1/2	0	0	0	0	0
0	1/3	0	0	0	0
0	0	1/3	0	0	0
0	0	0	1/4	0	0
0	0	0	0	1/5	

\tilde{D}^{-1}

New scale factor for columns

The diagram illustrates the computation of a feature vector X from the new adjacency matrix \tilde{A} and the new degree matrix \tilde{D}^{-1} . It shows two separate paths: one for the columns of \tilde{D}^{-1} and one for the columns of \tilde{A} .

Left path (columns of \tilde{D}^{-1}):

- Start with \tilde{D}^{-1} (green border). The column for node E is highlighted with a green border.
- Multiply by \tilde{A} (blue border). The row for node E is highlighted with a blue border.
- Result: Feature vector X (green border).

Right path (columns of \tilde{A}):

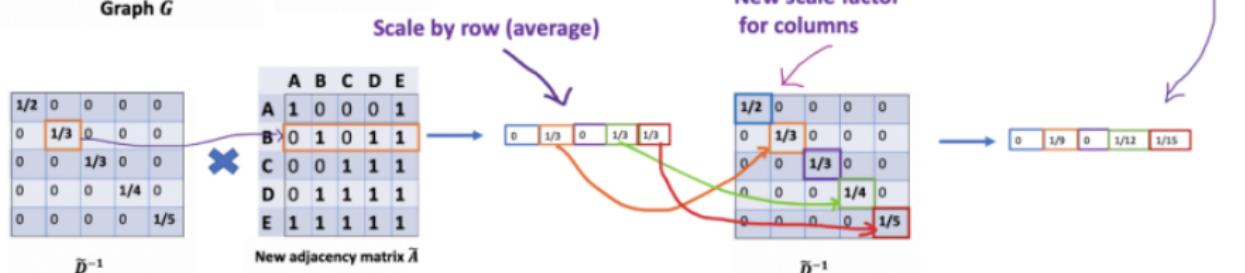
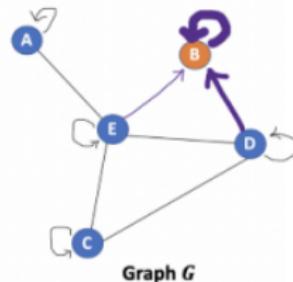
- Start with \tilde{A} (blue border). The column for node E is highlighted with a blue border.
- Multiply by \tilde{D}^{-1} (green border). The row for node E is highlighted with a green border.
- Result: Feature vector X (green border).

Feature vector X

A	-1.1	3.2	4.2
B	0.4	5.1	-1.2
C	1.2	1.3	2.1
D	1.4	-1.2	2.5
E	1.4	2.5	4.5

Neural Processes on Graphs

Related Work: Spectral Graph Convolutional Networks (GCN)



Neural Processes on Graphs

Related Work: Spectral Graph Convolution Networks (GCN)

2	0	0	0	0
0	3	0	0	0
0	0	3	0	0
0	0	0	4	0
0	0	0	0	5

\tilde{D}



1/2	0	0	0	0
0	1/3	0	0	0
0	0	1/3	0	0
0	0	0	1/4	0
0	0	0	0	1/5

\tilde{D}^{-1}

1/	0	0	0	0
0	1/	0	0	0
0	0	1/	0	0
0	0	0	1/2	0
0	0	0	0	1/

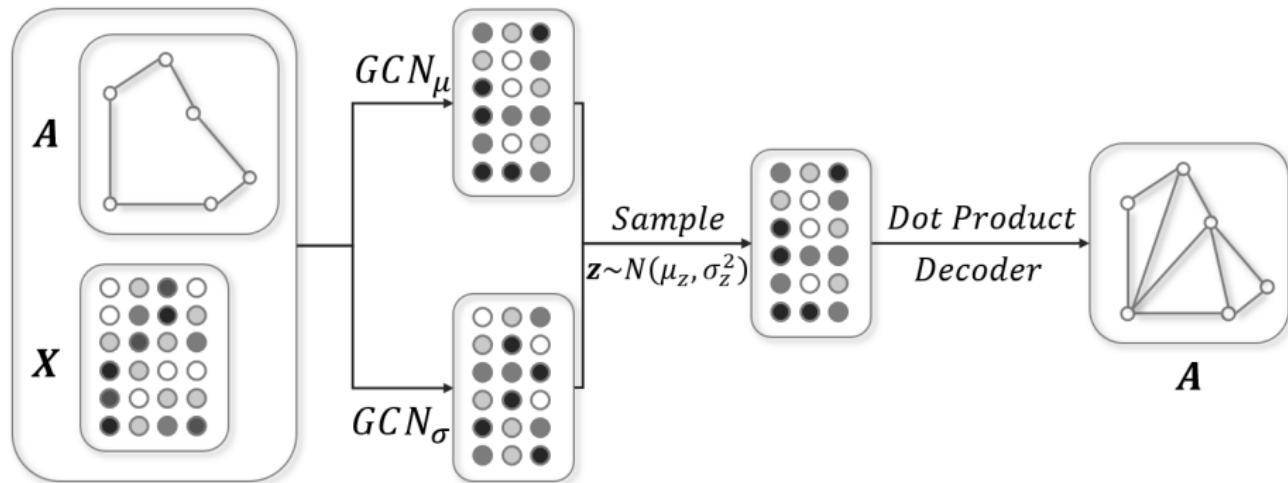
$\tilde{D}^{-1/2}$

$$Z = \text{ReLU}(\bar{A} \text{ReLU}(\bar{A} X W_1) W_2) \quad (23)$$

- Here $\bar{A} = \tilde{D}^{-\frac{1}{2}} \tilde{A} \tilde{D}^{-\frac{1}{2}}$, where $\tilde{D}_{ii} = \sum_j \tilde{A}_{ij}$ is the degree matrix of \tilde{A} , and $\tilde{A} = A + I$.

Neural Processes on Graphs

Related Work: Variational Graph Autoencoders Model



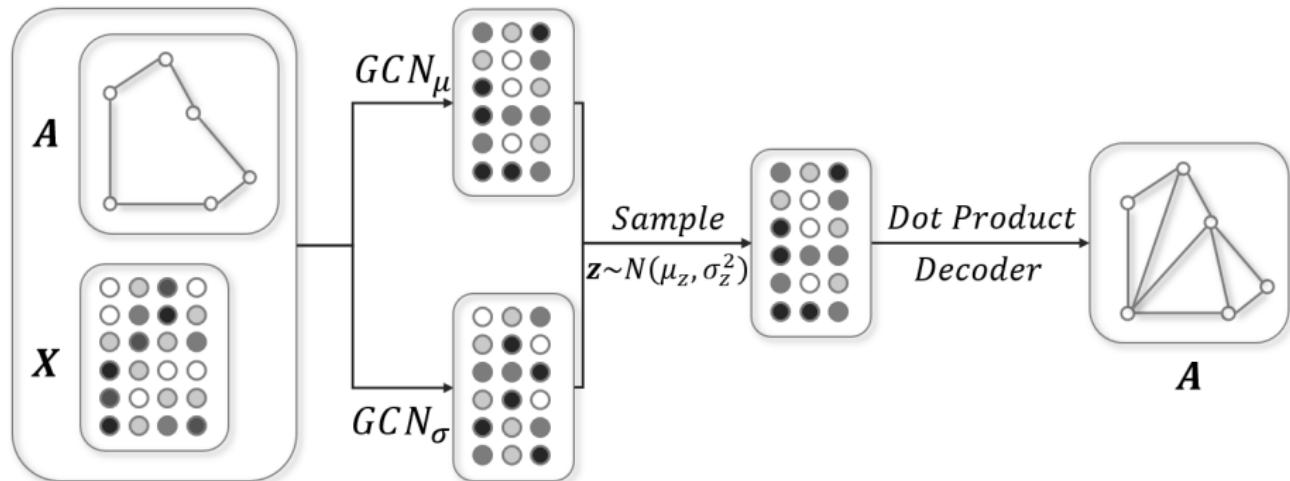
$$\boldsymbol{\mu} = \text{ReLU}(\bar{A} \text{ReLU}(\bar{A}XW_1)W_\mu) \quad (24)$$

$$\log \boldsymbol{\sigma} = \text{ReLU}(\bar{A} \text{ReLU}(\bar{A}XW_1)W_\sigma) \quad (25)$$

$$p(A|X, z) = \prod_{i=1}^n \prod_{j=1}^n p(A_{ij}|z_i, z_j), \quad \text{with} \quad p(A_{ij} = 1|z_i, z_j) = \sigma(z_i^\top z_j) \quad (26)$$

Neural Processes on Graphs

Related Work: Variational Graph Autoencoders Inference



$$\mathcal{L} = \mathbb{E}_{q(Z|A,X)}[\log p(A|Z)] - \mathbf{KL}[q(Z|A,X)||p(Z)] \quad (27)$$

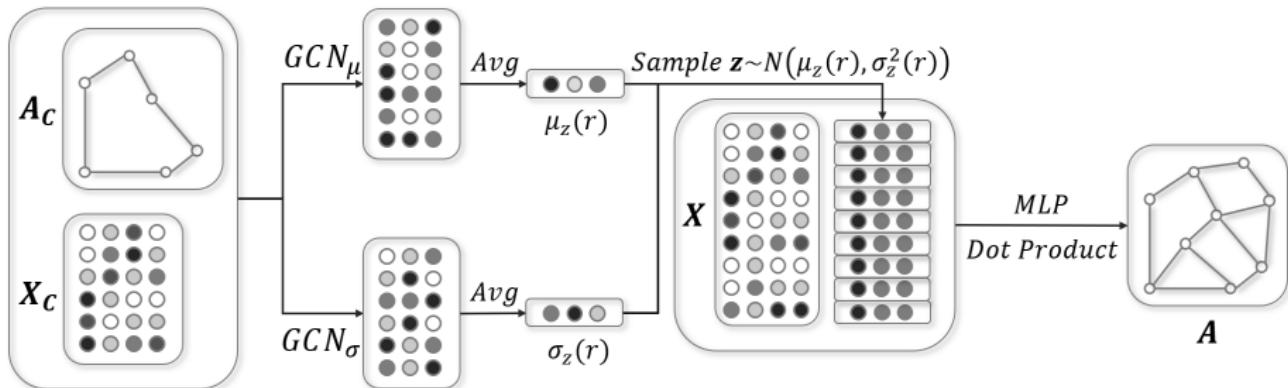
Neural Processes on Graphs

Graph Neural Processes Setup

- First we define an undirected graph $G = (V, A)$, where V is the set of nodes with features $x_i \in X$ corresponding to each node $v_i \in V$, and the adjacency matrix A with $A_{ij} = 1$ if there is a link between v_i and v_j , and $A_{ij} = 0$ otherwise.
- Then we randomly select a subset of nodes $V_C \in V$ with its related features $X_C \in X$ and adjacency matrix A_C to construct a context subgraph G_C . We assume there are n nodes in the complete graph G , and the first m nodes are the context nodes V_C .
- Our goal is to model the adjacency matrix A for the complete graph G conditional on the context subgraph G_C .

Neural Processes on Graphs

Graph Neural Processes Model Structure: Encoder



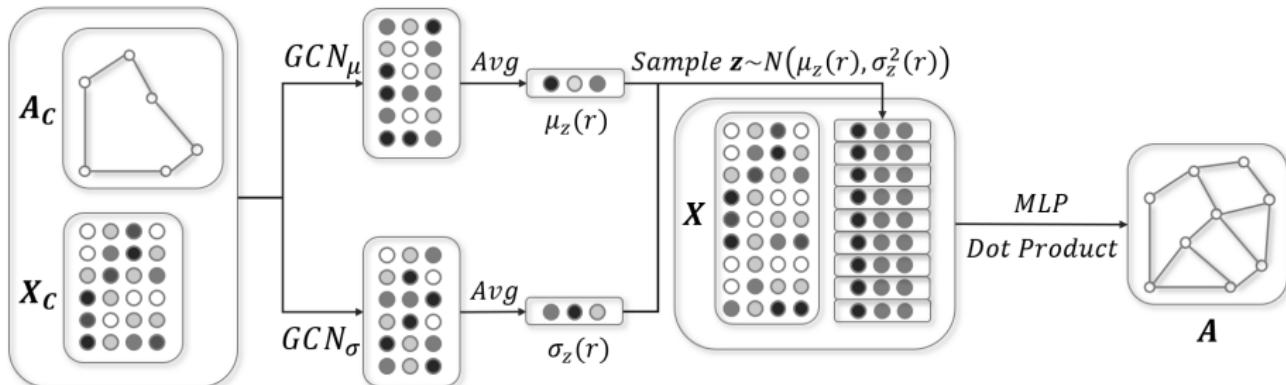
$$\mu_c = \text{ReLU}(\bar{A}_C \text{ReLU}(\bar{A}_C X_C W_1) W_\mu) \quad (28)$$

$$\log \sigma_c = \text{ReLU}(\bar{A}_C \text{ReLU}(\bar{A}_C X_C W_1) W_\sigma) \quad (29)$$

- Here $\bar{A}_C = \tilde{D}_C^{-\frac{1}{2}} \tilde{A}_C \tilde{D}_C^{-\frac{1}{2}}$, where $\tilde{D}_{C_{ii}} = \sum_j \tilde{A}_{C_{ij}}$ is the degree matrix of \tilde{A}_C , and $\tilde{A}_C = A_C + I$.

Neural Processes on Graphs

Graph Neural Processes Model Structure: Aggregation



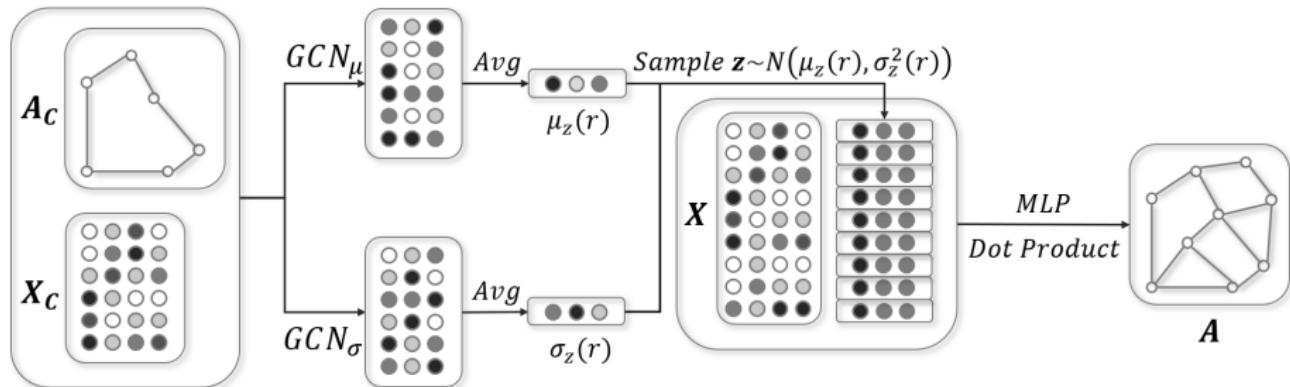
$$\mu_z(r) = \frac{1}{m} \sum_{i=1}^m \mu_{c_i} \quad \log \sigma_z(r) = \frac{1}{m} \sum_{i=1}^m \log \sigma_{c_i} \quad (30)$$

- Use $\mu_z(r)$ and $\sigma_z^2(r)$ to parameterise the distribution for the global latent representation.

$$q(z|A_C, X_C) = \mathcal{N}(\mu_z(r), \sigma_z^2(r))$$

Neural Processes on Graphs

Graph Neural Processes Model Structure: Decoder



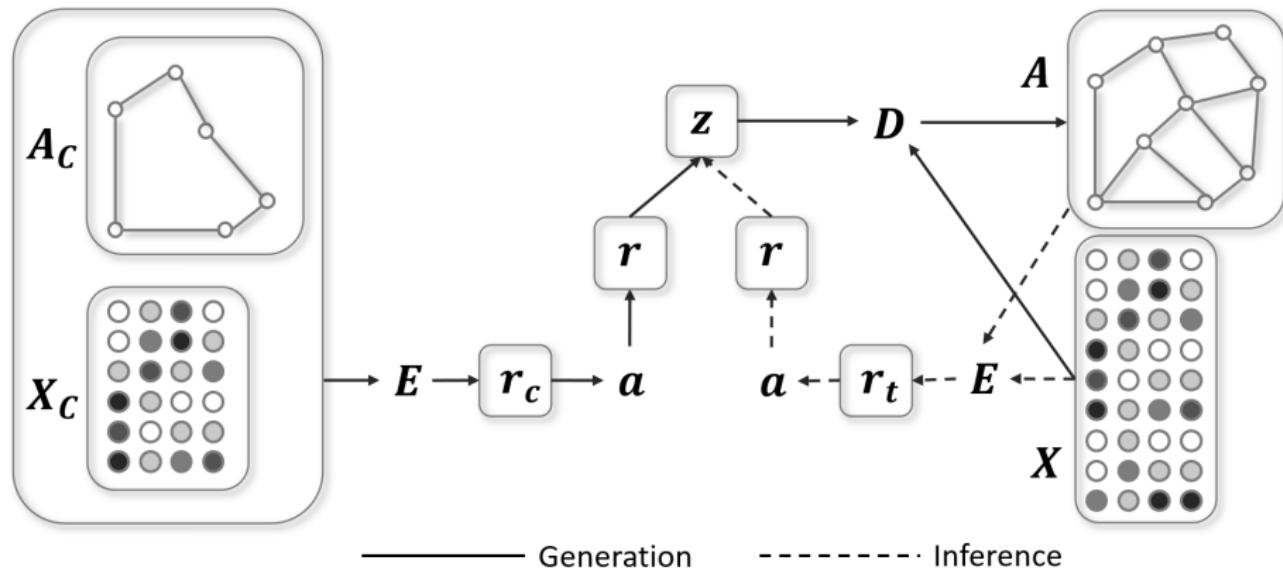
- Combine a sampled z with every feature $x_i \in X$ as $\tilde{x}_i \in \tilde{X}$, where $\tilde{x}_i = [x_i^\top \ z^\top]^\top$.

$$U = \sigma(W_3 \sigma(W_2 \tilde{X} + b_1) + b_2) \quad (31)$$

$$p(A|X, z) = \prod_{i=1}^n \prod_{j=1}^n p(A_{ij}|u_i, u_j), \quad \text{with } p(A_{ij} = 1|u_i, u_j) = \sigma(u_i^\top u_j) \quad (32)$$

Neural Processes on Graphs

Inference and Learning



$$\mathcal{L} = \mathbb{E}_{q(z|A, X)}[\log p(A|X, z)] - \mathbf{KL}[q(z|A, X)||q(z|A_C, X_C)] \quad (33)$$

The End